When Graph-Derived Neighbourhood Features Do Not Help: A Leakage-Controlled Evaluation of Behavioural-LRFM Churn Prediction in Non-Contractual Retail

FPT University, Ho Chi Minh City, Vietnam

Duc Nhat Vo[a], Gia Bao Duong[b]

[a]voducnhat242006@gmail.com, [b]ilyrabao@gmail.com

Abstract

Customer similarity graphs are an appealing addition to RFM-style churn models. If churn spreads through behavioural peer groups, the churn rate among a customer's nearest neighbours should help predict that customer's own churn. We test that idea on the Online Retail II dataset under two controls. First, the neighbourhood and community churn rates are recomputed inside every training fold, so no test label reaches the feature matrix. Second, every comparison runs over ten repetitions of five-fold cross-validation and is judged by the corrected resampled t-test rather than by a single split. Under these controls the idea does not hold. With the classifier fixed, no feature tier beats a plain three-variable RFM baseline, on either ROC-AUC or PR-AUC, for either classifier, at any graph degree from 5 to 20. The graph-enhanced tier gains +0.0033 ROC-AUC on one fold partition and reverses sign across ten. Every significant difference involves a change of classifier, so none speaks to the graph features. Attribution and contribution also come apart: the neighbourhood churn rate ranks third by SHAP, yet removing the graph features costs nothing. Because the k-NN graph is built on the features the classifier already uses, its neighbourhood churn rate largely restates existing information by construction; a graph built from co-purchase or referral links may behave differently. We argue that single-split evaluation cannot separate an effect this small from fold noise.

Keywords: Churn prediction · RFM · Customer similarity graph · Community detection · Repeated cross-validation · Negative result.

1 Introduction

Non-contractual retail gives no cancellation signal. A customer who stops buying does not announce it, so churn has to be defined by inactivity over a chosen horizon, and the modelling task is to predict that inactivity from purchase history. The RFM representation (recency, frequency, monetary value) has been the working vocabulary of this problem for decades, and a large applied literature extends it: adding relationship length, adding demographic attributes [9], or unrolling the triple over time [14].

A more recent line treats the customer base as a graph rather than a table. Corrêa Vianna Filho et al. formulate RFM segmentation as a maximum k-cut problem on a customer graph, where identical score vectors collapse into super-vertices and the reduced graph is bounded independently of the number of customers [6]. Once customers are vertices, a next step suggests itself: if a similarity graph can be built, then community detection [1, 17] can partition it, and graph-derived quantities such as degree, centrality and, in particular, the churn rate observed among a customer's neighbours can be appended to the feature matrix of a supervised churn model. The underlying hypothesis is a contagion one: churn risk is partly a property of the behavioural neighbourhood, not only of the individual.

This paper tests that hypothesis and does not find support for it. We build the pipeline the hypothesis calls for: a nine-dimensional behavioural extension of RFM, a k-nearest-neighbour customer similarity graph, Louvain and Leiden communities, and graph node features including neighbour and community churn rates. We then evaluate it against progressively simpler feature sets. The evaluation is where the contribution lies. Two controls turn out to be decisive:

1. **Fold-local recomputation of the neighbourhood features.** Both are aggregates of the target, so computing them once over the full dataset places test labels inside the training features. We recompute them inside every training fold from training labels only.
2. **Repeated cross-validation with a variance-corrected test.** A single split gives a point estimate with no uncertainty attached. We use ten repetitions on a shared fold partition and report the corrected resampled t-test of Nadeau and Bengio [15], which adjusts for the training-set overlap that makes the naive paired test anti-conservative [8, 2].

With both controls in place, the effect the pipeline was built to demonstrate disappears. We state our contributions conservatively:

- A negative result, reported as such: with the classifier fixed, neither the behavioural extension nor the graph features improve discrimination over a three-variable RFM baseline, at any graph degree tested.
- An evaluation protocol for this application area, together with the observation that differences below roughly 0.005 ROC-AUC on this dataset are distinguishable neither from fold noise nor from library-version drift.
- A dissociation between SHAP attribution rank and predictive contribution, which we state explicitly because feature-importance rankings are frequently read as evidence of usefulness.

2 Related Work

Three strands meet in this problem, and the tension between them is what makes a negative result informative rather than merely disappointing.

The *representation* strand asks what to measure about a customer. RFM persists because it is cheap and interpretable, and the extensions that follow it generally add dimensions rather than replace the idea: relationship length alongside the original three, demographic attributes [9], or a temporal unrolling in which the RFM triple becomes a sequence rather than a point [14]. What these extensions share is an implicit claim that the added dimensions carry information the original three do not. That claim is testable, and it is one of the things we test.

The *structural* strand asks whether customers should be modelled jointly rather than independently. Vianna Filho et al. make the strongest form of this argument for segmentation: by recasting RFM segmentation as maximum k-cut over a customer graph, the segmentation objective becomes an explicit graph quantity rather than a by-product of a distance-based clustering heuristic [6]. Their work is the starting point of ours, but it stops at segmentation and does not address supervised prediction. Community detection supplies the machinery for a supervised extension: modularity-based partitioning [16, 1] and its refinement guaranteeing well-connected communities [17] are standard, well-implemented, and apply directly to a similarity graph over customers. The step from there to appending the neighbourhood churn rate as a feature is short, which is precisely why it deserves scrutiny: that feature is an aggregate of the label, and features of this kind are the classic vehicle for target leakage.

The neighbourhood-churn feature is not itself new, and the settings where it has worked are informative. Telecommunications studies build it from call-detail records, where an edge means two subscribers actually spoke, and report that a subscriber's risk rises with the number of churned neighbours [11]; more recent work treats the social graph as a temporal object and predicts churn from how it evolves [12]. What separates those settings from ours is what an edge means. Theirs records an interaction that exists independently of the model's feature matrix; ours records proximity within the very feature space the classifier already sees. That difference turns out to matter, and Sect. 9.2 returns to it.

The *evaluation* strand is where the applied literature is weakest and where the methodological literature has long been settled. Overlapping resampled training sets inflate Type I error in the tests normally used to compare learning algorithms [8]; the variance correction that repairs the paired test for exactly this case has been available for two decades [15], along with evidence on how such tests replicate in practice [2, 7]. Applied churn studies, including our own earlier reporting of the present pipeline, nonetheless rank configurations by a single cross-validated point estimate. At differences on the order of 10^{-3} ROC-AUC that cannot separate an effect from a fold partition. Our results are an instance of the gap between these two strands, not a criticism of any individual study.

Gradient-boosted trees [5, 10] and additive feature attribution [13] are used here as standard tools, not as objects of study.

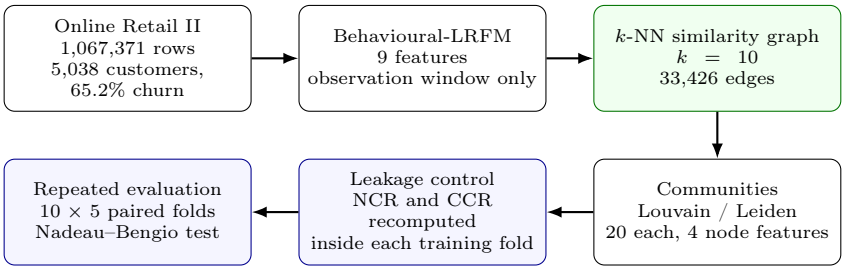


Fig. 1. Analysis pipeline. Green: the component under test. Blue: the two controls that determine the result, fold-local recomputation of the label-derived features and repeated evaluation with a variance-corrected test.

3 Pipeline Overview

Figure 1 shows the stages. The upper row builds the representation; the lower row, read right to left, carries it through the two controls into the evaluation.

4 Data and Feature Construction

4.1 Dataset and Windows

We use Online Retail II, the transaction log of a UK-based non-store retailer covering December 2009 to December 2011, distributed through the UCI Machine Learning Repository [3] and introduced in the RFM segmentation study of Chen et al. [4]. The raw file contains 1,067,371 transaction rows across two worksheets (525,461 and 541,910); the 541,000 figure often quoted for this dataset is the second worksheet alone.

Features and labels are separated in time. The observation window covers transactions from 2009-12-01 up to the timestamp 2011-06-30 00:00 and is the only period from which features are computed. The prediction window covers the following three months, up to 2011-09-30 00:00 (92 days); a customer active in the observation window is labelled a churner if no purchase appears in the prediction window. After filtering to valid transactions inside the observation window (556,212 rows) and aggregating to the customer level, the working matrix contains 5,038 customers. The churn rate is 65.2%.

That last number deserves emphasis, because it inverts an assumption often carried into this problem. Churners are the *majority* class here, not a rare positive to be rescued by resampling. A three-month inactivity horizon applied to a retailer with long interpurchase intervals labels most of the base as churned. This makes PR-AUC less informative than it would be under genuine imbalance, and we report it alongside ROC-AUC rather than as the primary metric.

4.2 Behavioural-LRFM Features

Nine customer-level features are computed inside the observation window: length L (days from the first purchase to the analysis date 2011-07-01), recency R (days from the last purchase to the same date), frequency F (distinct invoices), monetary value M, average order value, basket size, product diversity, cancellation rate, and mean interpurchase time. Cancellation rate is restricted to cancellations dated inside the observation window; an earlier version of this pipeline computed it over the full log, which allowed post-window cancellations to enter a pre-window feature.

5 Graph Construction and Community Detection

The nine features are standardised and a k-nearest-neighbour graph is built over customers with $k = 10$ and Euclidean distance, with edge weight $w(i, j) = 1/(1+d(i, j))$. The resulting graph has 5,038 nodes and 33,426 edges. Louvain [1] yields 20 communities with modularity 0.8385; Leiden [17] yields 20 communities with modularity 0.8323.

The high modularity should not be over-read: a k-NN graph with fixed k imposes near-uniform degree and strong local clustering by construction, so a value above 0.8 reflects how the graph was built at least as much as any structure in the customer base.

Six graph-derived features are appended: degree, weighted degree, PageRank, clustering coefficient, and the two neighbourhood aggregates that motivate the design: neighbour churn rate (mean label among a customer's own k nearest neighbours) and community churn rate (mean label within the customer's Louvain community).

6 Leakage Controls

Both neighbourhood aggregates are functions of the target. Computing either over the full dataset before cross-validation lets the label of a test customer influence the training features of its neighbours, and the resulting optimism can be substantial. We therefore recompute both inside each training fold: for customer i, the neighbour churn rate averages labels over $\mathcal{N}(i) \cap \mathcal{T}$, where $\mathcal{N}(i)$ is the set of the k nearest neighbours of i and $\mathcal{T}$ is the training index set, falling back to the training-fold base rate when that intersection is empty; the community churn rate is computed the same way within each Louvain community. The graph and the communities themselves are built from features only, never from labels, and are therefore fixed across folds.

Three further controls apply throughout. The classifier's scaler is fitted inside each training fold; the graph, its scaling and the K-means segments use all customers' features but no labels. K-means and Louvain cluster identifiers are one-hot encoded rather than passed as integers. Louvain and Leiden are both reported, but the repeated-CV features use Louvain.

7 Evaluation Protocol

7.1 Feature Tiers

Five nested tiers isolate the contribution of each modelling layer: **RFM** (R, F, M); **LRFM** (adding L); **B-LRFM** (the nine behavioural features); **B-LRFM+Seg** (adding 23 one-hot segment and community indicators); and **Graph-Enhanced**, the full 38-feature matrix adding the four static graph features and the two fold-local neighbourhood aggregates. Two classifiers are evaluated on every tier: LightGBM [10] as the primary model and ℓ_2-regularised logistic regression as an interpretable reference.

7.2 Repeated Cross-Validation and Significance

We run ten repetitions of stratified five-fold cross-validation, giving 50 folds. Within a repetition all five tiers and both classifiers see the *same* fold partition, so per-fold differences are properly paired. For a difference $d_1, \ldots, d_n$ over $n = 50$ folds we report the corrected resampled t-statistic [15]

$$t = \frac{\bar{d}}{\sqrt{\hat{\sigma}^2 \left(\frac{1}{n} + \frac{n_2}{n_1} \right)}}, \qquad \frac{n_2}{n_1} = \frac{1}{4}, \tag{1}$$

where n_2/n_1 is the ratio of test-set to training-set size, which for five-fold cross-validation is $1/4$. We report p_{NB} from this statistic only. Bootstrap confidence intervals over the 50 differences are given for effect magnitude, but we caution against reading them as significance statements: the bootstrap resamples the differences as if independent, while repeated folds share training data, so those intervals are narrower than the correlation structure warrants. This is the same overlap that makes the naive paired t-test anti-conservative [8].

Run with the single seed used by the original pipeline, the significance harness reproduces that pipeline's per-tier ROC-AUC to four decimal places on all five tiers, so any divergence below is attributable to the fold partition rather than to a re-implementation difference.

8 Results

8.1 One Split Says the Opposite of Ten

Table 1 puts the two evaluations side by side, and the contrast is the result of this paper. On the single stratified five-fold split this pipeline originally reported, Graph-Enhanced is the best tier, +0.0033 ROC-AUC above RFM. Over ten repetitions it is jointly the worst, −0.0021 below RFM. Nothing about the model changed between those two columns; only the fold partition did. Even at the seed that produces it, the +0.0033 gap has a bootstrap 95% interval of $[-0.0025, +0.0099]$, which contains zero.

Table 1. LightGBM across the five feature tiers. The single-split column is the seed-42 evaluation this pipeline first reported; the remaining columns are ten repetitions of the same protocol. p_{NB} is the corrected resampled t-test against RFM (Sect. 7.2); no tier reaches significance. Why the single-split column misleads is taken up in Sect. 9.

	Single split	Repeated CV (10×5 folds, paired)		
Feature tier	ROC-AUC	ROC-AUC	Δ vs RFM	p_{NB}
RFM (baseline)	0.8065	0.8066		
LRFM	0.8071	0.8074	+0.0009	0.755
B-LRFM	0.8067	0.8054	−0.0012	0.742
B-LRFM+Seg	0.8073	0.8044	−0.0022	0.561
Graph-Enhanced	**0.8098**	0.8044	−0.0021	0.590

Table 2. Classifier comparisons on the same 50 paired folds. Bootstrap intervals show magnitude only; significance is read from p_{NB} (Sect. 7.2).

Comparison	Δ ROC-AUC	95% CI	p_{NB}
(a) Classifier and tier both change			
LogReg / B-LRFM vs. LGBM / Graph-Enh.	+0.0113	[+0.0096, +0.0130]	**0.001**
LogReg / RFM vs. LGBM / Graph-Enh.	+0.0099	[+0.0080, +0.0119]	**0.010**
(b) Same tier, classifier only			
LogReg vs. LGBM, both on RFM	+0.0078	[+0.0056, +0.0100]	0.068
LogReg vs. LGBM, both on Graph-Enh.	+0.0108	[+0.0091, +0.0126]	**0.002**

No tier differs significantly from the RFM baseline, and three of the four point differences are negative. The same holds for PR-AUC and for logistic regression as the classifier: across the four metric-classifier combinations (sixteen tier-versus-baseline comparisons), the smallest p_{NB} is 0.144. On the single split the other metrics agree with ROC-AUC, with Graph-Enhanced leading on PR-AUC (0.8689), F1 (0.8108) and recall (0.8009); that agreement is what makes a single split persuasive and is exactly why it is not sufficient.

8.2 Classifier Effects Are Not Feature Effects

Table 2 compares the two classifiers. In group (a), logistic regression on the three RFM variables beats the 38-feature gradient-boosted graph model ($p_{\mathrm{NB}} = 0.010$). These rows change the classifier and the feature tier together, so they cannot tell a model effect from a feature effect. Group (b) keeps the tier fixed. Logistic regression is ahead at every tier, by +0.0062 to +0.0108 ROC-AUC. The gap is significant at the three tiers with nine or more features ($p_{\mathrm{NB}} \leq 0.017$) but not on RFM (0.068) or LRFM (0.154), and only the Graph-Enhanced row survives a Bonferroni correction over the five tiers.

Conclusions about features therefore rest on Table 1 and its logistic-regression counterpart, where the classifier is fixed and no tier differs from RFM (Sect. 8.1).

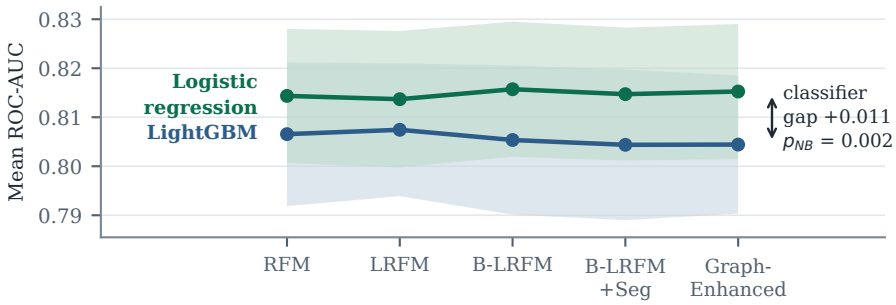


Fig. 2. Mean ROC-AUC over 50 folds for each tier and classifier. Bands show ±1 standard deviation of the fold scores, the spread a single split is drawn from.

Table 3. Graph-Enhanced versus the RFM baseline as the graph degree varies. LightGBM, 10×5 paired folds at every setting; the graph, the Louvain communities and the fold-local features are rebuilt for each k.

k	Edges	Communities	Δ vs RFM	p_{NB}
5	16,902	31	−0.0006	0.880
10	33,426	20	−0.0021	0.590
15	50,054	16	+0.0017	0.666
20	66,652	16	−0.0027	0.500

This fits RFM already carrying most of the signal on this dataset, which we do not test directly.

Figure 2 shows both effects. Along each line, adding tenure, behaviour, segments and then the graph changes nothing significant. The vertical gap between the lines is the classifier effect of Table 2(b). The shaded bands show why one split cannot settle this: fold-to-fold spread is about 0.014 standard deviation, an order of magnitude larger than any difference between tiers.

8.3 The Result Does Not Depend on the Graph Degree

A negative result about a graph may be a result about one particular graph. We therefore rebuilt the graph, the communities and the fold-local features at $k \in \{5, 10, 15, 20\}$ and reran the full repeated-CV protocol at each setting. Table 3 reports the outcome.

Quadrupling the edge count and halving the number of communities changes the graph substantially and the conclusion not at all: no setting reaches significance, and the sign of the difference is not even stable, turning positive at $k = 15$ and negative at the three others. A quantity whose sign depends on a tuning parameter it should not depend on is behaving like noise, which is what the p-values say directly.

8.4 Attribution Is Not Contribution

Global SHAP values [13] on the graph-enhanced model rank recency first (1.087) and monetary value second (0.560), followed by neighbour churn rate (0.250), interpurchase time (0.224) and community churn rate (0.215). Read as an importance ranking, this appears to confirm that the neighbourhood aggregates matter: two of the top five features are graph-derived.

They do not matter for discrimination: the tier without the six graph features (B-LRFM+Seg, 0.8044) and the tier with them (Graph-Enhanced, 0.8044) are indistinguishable. This ablation drops all six graph features at once and does not isolate either aggregate. This is a matter of definition rather than a contradiction. SHAP attributes how a *fitted* model distributes credit across the features it was given, so a feature correlated with both the target and the rest of the feature set absorbs attribution while adding nothing the others did not already carry. Attribution rank answers "what is this model using?"; only an ablation answers "would removing it cost anything?", and here it answers no.

8.5 Relation to the Segmentation Baseline

Our starting point [6] segments customers on discrete 1–5 RFM scores in three dimensions; we predict churn in a nine-dimensional continuous space. Silhouette scores are not comparable across feature spaces, so we make no clustering-quality claim: the baseline reports 0.398 at $k = 4$ on its own scoring space, while K-means and Louvain on our space reach 0.219 and -0.050. The relation between the two works is scope extension, not improvement. The communities remain usable for prioritisation: ranked by churn probability times monetary value, they separate a high-value at-risk group (327 customers, 82% churn) from a high-value loyal one (418 customers, 24% churn), a descriptive output only, since the data contain no campaign responses.

9 Discussion

9.1 Effect Sizes Below the Noise Floor

Two independent sources of variation are the same size as the effects the pipeline was built to detect. The fold partition contributes roughly 0.003 ROC-AUC, as the sign reversal above shows. Library versions contribute as much: re-running the auxiliary improvement pipeline under newer libraries moved its best configuration by -0.0033 ROC-AUC and the Leiden community count from 20 to 21. We therefore treat differences below roughly 0.005 ROC-AUC on this dataset as uninformative, and suggest that comparisons here either report a variance-corrected test or refrain from ranking configurations inside that band.

9.2 Why the Contagion Hypothesis Might Still Hold

Our result is a failure to reject a null on one dataset with one graph construction, not a demonstration that behavioural peer effects are absent. The most likely explanation is that we built the wrong graph: its edges encode proximity in a feature space derived from the same purchase history that already feeds the classifier, so the neighbour churn rate is plausibly a smoothed restatement of what the model has rather than new information. This reading is consistent with where the feature does work. The telecommunications and social graph results cited in Sect. 2 are built on recorded interaction (who called whom, who plays with whom), not on similarity computed from the model's own inputs [11, 12].

Churn itself also differs between the two settings. In telecommunications it is the cancellation of a service contract, and peers have a way to influence each other: people who call each other benefit from staying on the same network, and a subscriber's churn risk has been linked to the number of contacts who have already left [11]. In non-contractual retail nobody cancels. Churn is inferred from a gap in purchases, and two customers who buy alike need not know each other. A similarity graph therefore connects customers who resemble each other, not customers who influence each other, and its neighbourhood churn rate can at most summarise what similar customers did, which the classifier already sees in the features. A retail graph built from interactions absent from the feature matrix, such as shared baskets, referral or household links, would be a different test, and is the experiment we would run next. Two weaker readings also fit: a three-month inactivity label on a majority-churn base may be too coarse for neighbourhood effects to surface, and with all tiers near 0.81 ROC-AUC the binding constraint may be the representation rather than the model class.

10 Threats to Validity

All results come from one retailer over two years, so we make no claim about other sectors or horizons, and the effect-size floor we report holds only for this dataset and sample size. The degree sweep rules out an artefact of k, but every graph we tried is a k-NN graph on the same features, and graphs built from interactions remain untested (Sect. 9.2). Two parts of the original study design, a country-of-origin encoding and spectral clustering, were never implemented and play no role in the results. Because churners are the majority class, PR-AUC loses its usual meaning here, and no conclusion rests on it alone.

Our literature search was targeted rather than systematic. The studies of neighbourhood churn that we found (Sect. 2) use interaction graphs, but we make no priority claim. Finally, the expected retention value ranks customers by a model output, not by observed campaign response, so its practical value is untested.

11 Conclusion

We built the pipeline the behavioural-peer-effect hypothesis calls for and evaluated it under fold-local recomputation of the label-derived features and repeated cross-validation with a corrected test. The hypothesis is not supported on this dataset at any graph degree tested. The significant differences we found are between classifiers and say nothing about the graph features.

The methodological point generalises further than the negative result. An evaluation that reports one cross-validated number per configuration cannot distinguish a 0.003 ROC-AUC difference from the partition that produced it. The controls used here are inexpensive (ten repetitions of five-fold cross-validation on 5,038 customers take minutes), and we would argue they belong in the default protocol rather than in an appendix.

Reproducibility. Code, per-fold scores, significance summaries and raw run logs are available at `https://github.com/VoDucNhatku/graph-behavioral-lrfm-churn`, tagged `v1.1`. Library versions are pinned to the environment that produced every reported number, and Fig. 2 and Table 2 are computed from the committed per-fold scores. The full pipeline runs in under ten minutes on a laptop with no GPU.

Acknowledgments. The authors used a large language model (Claude, Anthropic) to draft and revise the manuscript, write parts of the evaluation and figure code, and cross-check numbers against the result files. All numbers were verified against the archived runs, and the authors take full responsibility for the content.

Disclosure of Interests. The authors have no competing interests to declare that are relevant to the content of this article.

References

1. Blondel, V.D., Guillaume, J.L., Lambiotte, R., Lefebvre, E.: Fast unfolding of communities in large networks. Journal of Statistical Mechanics: Theory and Experiment **2008**(10), P10008 (2008). https://doi.org/10.1088/1742-5468/2008/10/P10008
2. Bouckaert, R.R., Frank, E.: Evaluating the replicability of significance tests for comparing learning algorithms. In: Advances in Knowledge Discovery and Data Mining (PAKDD 2004). Lecture Notes in Computer Science, vol. 3056, pp. 3–12 (2004). https://doi.org/10.1007/978-3-540-24775-3_3
3. Chen, D.: Online retail II. UCI Machine Learning Repository (2019). https://doi.org/10.24432/C5CG6D, `https://archive.ics.uci.edu/dataset/502/online+retail+ii`
4. Chen, D., Sain, S.L., Guo, K.: Data mining for the online retail industry: A case study of RFM model-based customer segmentation using data mining. Journal of Database Marketing & Customer Strategy Management **19**(3), 197–208 (2012). https://doi.org/10.1057/dbm.2012.17

5. Chen, T., Guestrin, C.: XGBoost: A scalable tree boosting system. In: Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining. pp. 785–794 (2016). https://doi.org/10.1145/2939672.2939785
6. Corrêa Vianna Filho, A.L., de Lima, L., Kleina, M.: RFM model customer segmentation from a graph theory perspective. Quality & Quantity **60**(5), 15365–15386 (2026). https://doi.org/10.1007/s11135-026-02784-0
7. Demšar, J.: Statistical comparisons of classifiers over multiple data sets. Journal of Machine Learning Research **7**, 1–30 (2006), https://www.jmlr.org/papers/v7/demsar06a.html
8. Dietterich, T.G.: Approximate statistical tests for comparing supervised classification learning algorithms. Neural Computation **10**(7), 1895–1923 (1998). https://doi.org/10.1162/089976698300017197
9. Ho, T., Nguyen, S., Nguyen, H., Nguyen, N., Man, D.S., Le, T.G.: An extended RFM model for customer behaviour and demographic analysis in retail industry. Business Systems Research Journal **14**(1), 26–53 (2023). https://doi.org/10.2478/bsrj-2023-0002
10. Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., Ye, Q., Liu, T.Y.: LightGBM: A highly efficient gradient boosting decision tree. In: Advances in Neural Information Processing Systems 30 (NIPS 2017). pp. 3146–3154 (2017)
11. Kostić, S.M., Simić, M.I., Kostić, M.V.: Social network analysis and churn prediction in telecommunications using graph theory. Entropy **22**(7), 753 (2020). https://doi.org/10.3390/e22070753
12. Lee, M., Woo, J.: TempODEGraphNet: predicting user churn using dynamic social graphs and neural ODEs. PLOS One **20**, e0321560 (2025). https://doi.org/10.1371/journal.pone.0321560
13. Lundberg, S.M., Lee, S.I.: A unified approach to interpreting model predictions. In: Advances in Neural Information Processing Systems 30 (NIPS 2017). pp. 4765–4774 (2017)
14. Mena, G., Coussement, K., De Bock, K.W., De Caigny, A., Lessmann, S.: Exploiting time-varying RFM measures for customer churn prediction with deep neural networks. Annals of Operations Research **339**(1–2), 765–787 (2023). https://doi.org/10.1007/s10479-023-05259-9
15. Nadeau, C., Bengio, Y.: Inference for the generalization error. Machine Learning **52**(3), 239–281 (2003). https://doi.org/10.1023/A:1024068626366
16. Newman, M.E.J., Girvan, M.: Finding and evaluating community structure in networks. Physical Review E **69**(2), 026113 (2004). https://doi.org/10.1103/PhysRevE.69.026113
17. Traag, V.A., Waltman, L., van Eck, N.J.: From Louvain to Leiden: guaranteeing well-connected communities. Scientific Reports **9**, 5233 (2019). https://doi.org/10.1038/s41598-019-41695-z